


Deep learning




Mannu Malhotra

Until Now:



Now:





Agenda

- Neural Networks 101:
 - What are NN
 - Why do we need NN
 - What are computation graphs
 - Derivation on computation graphs
 - **Hands on: Implement logistic regression in NN style**
 - NN representation
- Advanced topics:
 - Activation functions
 - Gradient descent and Backpropagation
 - Initialization of weights
- Practical Approach
 - Building blocks of DNN, forward and back prop
 - Softmax and **hands on: Softmax regression using Keras**
 - Explain homework: Building a NN in python using numpy.



What are Neural Networks

- We learned linear regression
 - $y = wx + b$
- Logistic regression : linear regression with sigmoid
 - For binary classification.
 - Uses sigmoid activation.
 - Get the value between 0 and 1 to determine the class.
- Logistic regression : single layer, single unit neural network.



Why do we need Neural Networks?

- Linearly non-separable data
 - Linear models like logistic regression can not learn on such data.

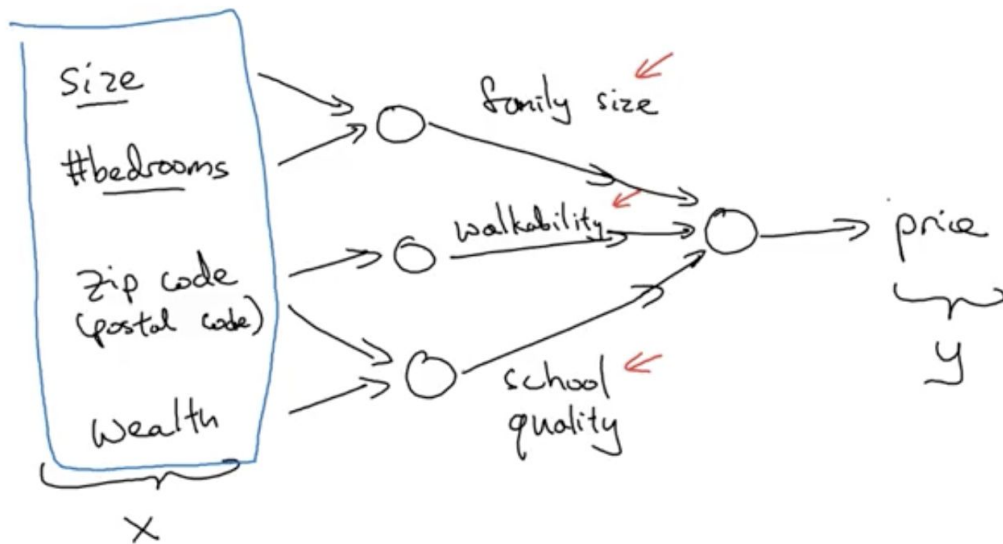


- Use polynomial features
- Use kernels to do features transformation to new dimensions.
- Use stacking of logistic regression to build non linear learner.
 - NN : $y = w_2 \cdot \text{sigmoid}(w_1 \cdot x + b_1) + b_2$

Why do we need Neural Networks?


- Feature learning/end-to-end learning!

Housing Price Prediction




Computation graphs & chain rule for first derivative

- A NN is nothing but a computation graph.
- Computation of neural network = forward propagation (output of the neural network) + backward propagation (for optimization like gradient descent).
- Illustration on chain rule : how would I calculate dJ/da (viz. da).



Logistic regression derivatives



$$z^{(i)} = w^T x^{(i)} + b$$

$$\hat{y}^{(i)} = a^{(i)} = \text{sigmoid}(z^{(i)})$$

$$\mathcal{L}(a^{(i)}, y^{(i)}) = -y^{(i)} \log(a^{(i)}) - (1 - y^{(i)}) \log(1 - a^{(i)})$$


- Objective: Modify w & b to $\min(L)$
- Calculate derivative of Loss function w.r.t. to w and b .
 - i.e. calculate dw & db
- $da = -y/a + (1-y)/(1-a)$
- $da/dz = a(1-a)$
- $dz = a - y$
- $db = dz, dw = xdz$

Gradient descent: One last time!

Recap: $\hat{y} = \sigma(w^T x + b)$, $\sigma(z) = \frac{1}{1+e^{-z}}$ ←

$$\underline{J(w, b)} = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})$$


- Objective: Find w, b to minimize J (cost).





Gradient descent algorithm

repeat until convergence {
 $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$
 (for $j = 1$ and $j = 0$)
}

Hands on: Implement logistic regression in NN style




NN representation : logistic regression vs NN




NN representation


- Forward prop:
 - cal $z^{[1]}, a^{[1]}$; $w^{[2]} = a^{[1]}$; cal $z^{[2]}, a^{[2]}$; pred = $a^{[2]}$; cal $L(\text{pred}, y)$
- Backward prop:
 - Cal $da^{[2]}, dz^{[2]}, dw^{[2]}, db^{[2]}, da^{[1]}, dz^{[1]}, dw^{[1]}, db^{[1]}$



NN representation : layers



NN representation :



$$z_1^{[1]} = w_1^{[1]T} x + b_1^{[1]}, \quad a_1^{[1]} = \sigma(z_1^{[1]})$$

$$z_2^{[1]} = w_2^{[1]T} x + b_2^{[1]}, \quad a_2^{[1]} = \sigma(z_2^{[1]})$$

$$z_3^{[1]} = w_3^{[1]T} x + b_3^{[1]}, \quad a_3^{[1]} = \sigma(z_3^{[1]})$$

$$z_4^{[1]} = w_4^{[1]T} x + b_4^{[1]}, \quad a_4^{[1]} = \sigma(z_4^{[1]})$$

Back prop:

$$z^{[1]} = w^{[1]} X + b^{[1]}$$

$$a^{[1]} = g^{[1]}(z^{[1]})$$

$$z^{[2]} = w^{[2]} A^{[1]} + b^{[2]}$$

$$A^{[2]} = g^{[2]}(z^{[2]})$$

$$dz^{[2]} = A^{[2]} - Y$$


$$dw^{[2]} = dz^{[2]} * A^{[1]} / m \quad (X \rightarrow A^{[1]})$$

$$db^{[2]} = \sum dz^{[2]} / m$$

$$dz^{[1]} = w^{[2]} * dz^{[2]} * g^{[1]'}(z^{[1]})$$


$$dw^{[1]} = dz^{[1]} X^T$$

$$db^{[1]} = \sum dz^{[1]} / m$$




Activation functions: Which Activation functions to use?


1. Linear:



3. Sigmoid:



2. Relu:



4. Softmax




Activation functions

- Default choice for hidden units.
- Prevent from vanishing gradients.





Activation functions: Why do we need activation functions?

- Why not linear activation function(no activation)?
- ~~$a = g(z)$~~ , $a = z$.
- Y_{pred} = linear function to input parameters.
- $z_1 = w_1.x + b$
- $a_1 = g(z_1) = z_1$
- $a_2 = w(a_1)+b = w_2.z_1 + b_2 = w_2(w_1.x + b_1) + b_2$
- $Y_{\text{pred}} = A_2 = w_2.w_1.x + w_2.b_1 + b_2$



Initialization of weights

- All hidden units learning the same function : symmetric breaking.
- Initialise with small values of weights to make learning faster:
 - Large value of weights ->
 - High value of z ->
 - Flat part for sigmoid
 - Tangent value to be small.



Multiclass classification? : Softmax

- What if we have multiple classes in the target?
- We can have:
 - One vs all (One vs rest)
 - # of models trained: N (Number of classes)
 - Problems: Class imbalance, can not capture the interdependencies of the classes.
 - One vs One (One vs One)
 - # of models trained: $N(N-1)/2$
 - Problem: High time complexity, can not capture the interdependencies of the classes.
- Softmax layer to the rescue!


Multiclass classification? : Softmax

- N units in the output layer.
- Sigmoid vs Softmax:
 - Just a different activation function.

$$\frac{1}{1+e^{-(w^T x+b)}}$$

$$\frac{e^{y_i}}{\sum_j e^{y_j}}$$

- Just a normalization of results.



Which deep learning framework to use?




Keras



PYTORCH



Caffe



Chainer

Which deep learning framework to use?



He said: "I will just install Caffe and I will come back"


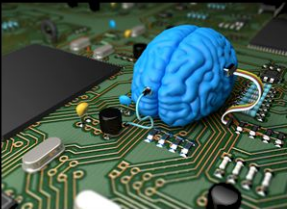





I'm still waiting for him

Hands on: Softmax regression using Keras

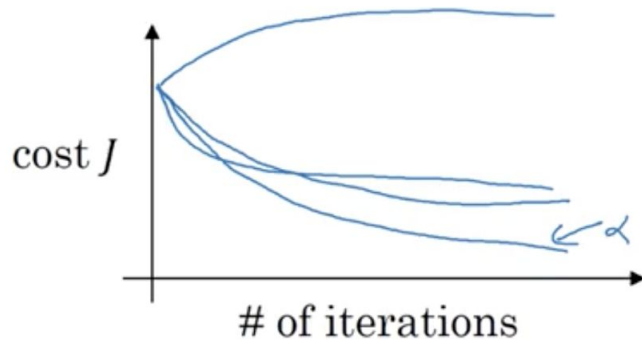

-

Deep Learning


 <p>What society thinks I do</p>	 <p>What my friends think I do</p>	 <p>What other computer scientists think I do</p>
 <p>What mathematicians think I do</p>	 <p>What I think I do</p>	<pre>In [1]: import keras Using TensorFlow backend.</pre> <p>What I actually do</p>

Hyperparameters!

- Parameters which can not be learned.
 - E.g. learning rate, number of layers, number of units, epochs (iterations), choice of activation functions.
- How to find the optimal values?
 - Trial and error and plot the cost function.



Hyperparameters!



Terima kasih!



Explain homework!

