

TREES

-Mannu (`_mannuscript`)

WHAT ARE TREE BASED ALGORITHMS?

- Set of:
 - Non-linear,
 - Supervised algorithms
 - For classification and regression.
- Examples:
 - Decision trees,
 - Random forests,
 - Gradient boosting.

WHY TREES?



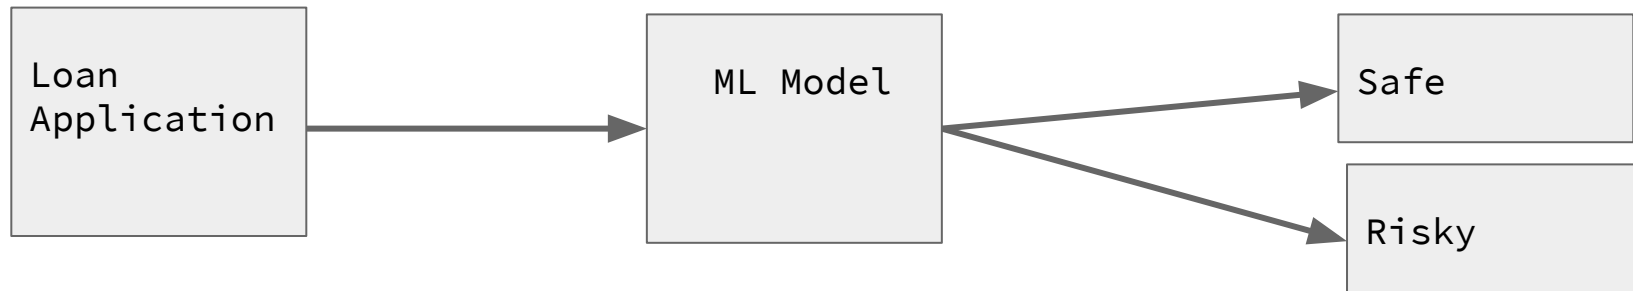
WHY TREES?

- Easy to understand.
- Very less fancy statistics required.
- Not data hungry.
- Not affected by outliers.

DECISION TREES

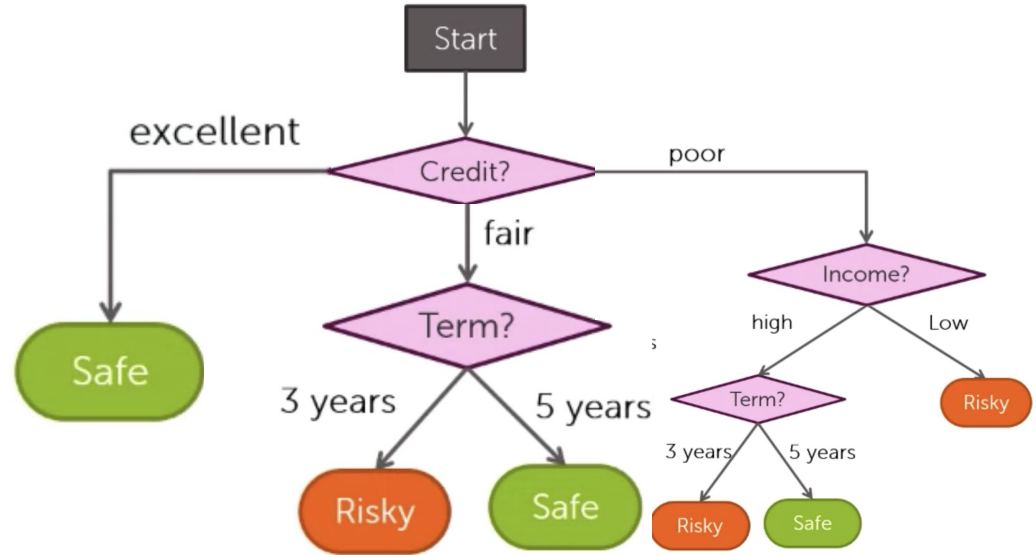
PROBLEM

- Predict if a given loan application is safe or risky.
- Given the data:
 - Credit History
 - Income
 - Term
 - Personal Information



PROBLEM

- Credit History
- Income
- Term
- Personal Information




DECISION TREES

- Objective:

- Given a dataset, create a Tree $T(x)$ with minimum classification error on the training dataset.
- Classification error = # of incorrect prediction / # of samples.



- But wait!

- How to decide the feature/node placements?




DECISION TREES

- Objective: find t from T to $\min(E)$.
 - T : Set of all possible decision tree for give features.
 - E : Classification error on training data.
 - t : Best tree.
- Easy: Just calculate E for all t in T !



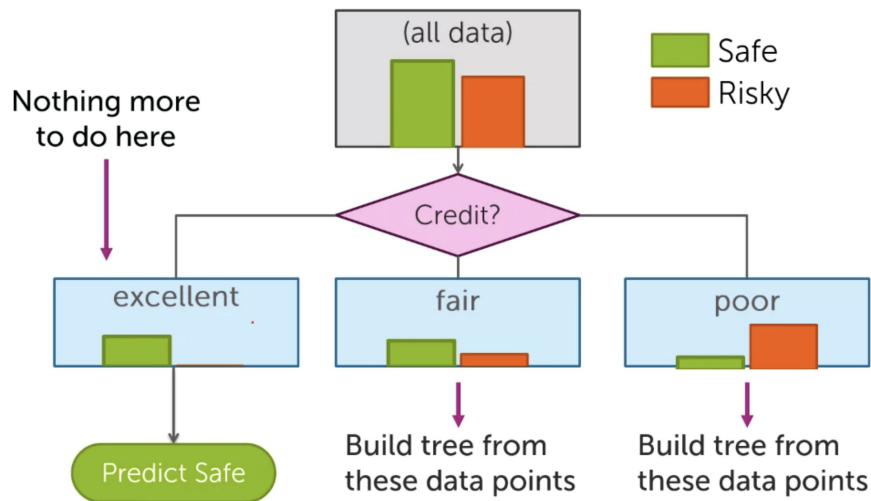
DECISION TREES

- There is a problem:
 - NP-hard problem!!!
 - # of trees grows exponentially.
 - No approximate solution.
- Need some heuristic to build the tree:
 - Greedy algorithm.



DECISION TREES: GREEDY ALGORITHM

1. Start with an empty tree.
2. Select a feature to split.
3. For each split:
 - a. If nothing more, stop the recursion.
 - b. Otherwise, recurse, go back to step 2 and continue to split.



DECISION TREES: GREEDY ALGORITHM

1. Start with an empty tree.
 2. Select a feature to split.
 3. For each split:
 - a. If nothing more, stop the recursion.
 - b. Otherwise, recurse, go back to step 2 and continue to split.
- Two challenges:
 - How to decide which feature to split at step 2?
 - What are the stopping conditions at step 3.a?

DECISION TREES

- What are the stopping conditions?
 - Nothing to split.
 - No more features left.
 - Constraint on max depth of the tree.

DECISION TREES

- How to decide which feature to split first?
 - Gini Index
 - Information Gain
 - Classification Error

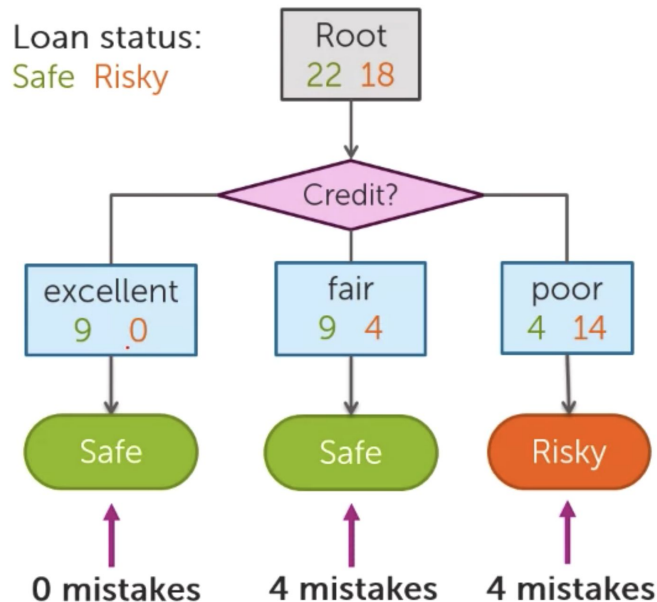
DECISION TREES: FEATURE SPLIT

- Classification Error:
 - a. For each (remaining) feature:
 - Split on the values
 - For each split:
 - Assign majority class as the prediction class
 - Calculate the misclassified examples M_s
 - Calculate $E^f = (\sum_s M^f) / \text{Total examples}$
 - b. Choose the feature with min E^f .

DECISION TREES: FEATURE SPLIT

- Classification Error:
 - a. For each (remaining) feature:
(Choose Credit)
 - Split on the values (**excellent**, **fair**, **poor**)
 - For each split:
 - Assign majority class as the prediction class.
(Safe, Safe, Risky)
 - Calculate the misclassified examples M_s (0, 4, 4)
 - Calculate E^f (0 + 4 + 4)/40


Choice 1: Split on Credit



DECISION TREES: FEATURE SPLIT

- Classification Error:
 - a. For each (remaining) feature: (**Choose Term**)
 - Split on the values (**3y, 5y**)
 - For each split:
 - Assign majority class as the prediction class. (**Safe, Risky**)
 - Calculate the misclassified examples M_s (**4, 6**)
 - Calculate $E^f = (4 + 6)/40$

Choice 2: Split on Term



DECISION TREES: FEATURE SPLIT

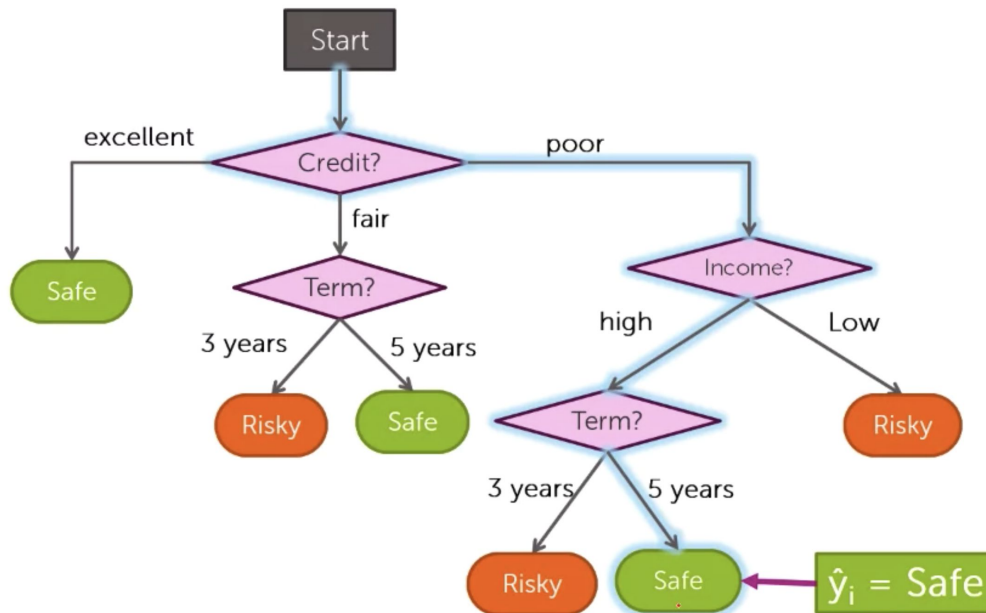
- Classification error:

- Credit: $(0 + 4 + 4)/40 = .2$
- Term: $(4 + 6)/40 = .25$

1. Start with an empty tree.
2. Select a feature to split.
3. For each split:
 - a. If nothing more, stop the recursion.
 - b. Otherwise, recurse, go back to step 2 and continue to split.

DECISION TREES: PREDICTING VALUES


- Predicting = Traversing the tree
- E.g. (Credit = Poor, Term = 5y, Income = High)



DECISION TREES: CONTINUOUS VALUED FEATURES


- Features are not only categorical.
- How to split now?
 - a. Number of classes are infinite.
- Use threshold split.
- Algorithm:
 - a. Sort the values
 - b. For $i = 1$ to $N-1$:
 - $t_i = (N_i + N_{i+1})/2$
 - Compute classification error
- Choose the N_i with $\min(\text{error})$ as the threshold split.

DECISION TREES: HANDS ON



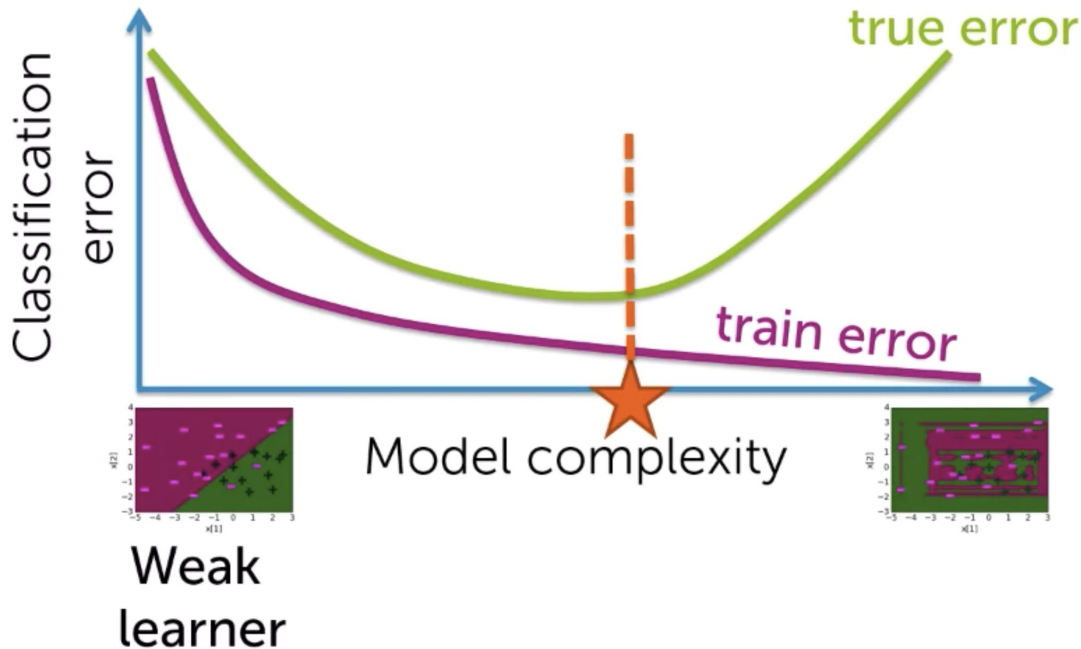
OVERFITTING IN DECISION TREES

- Early stopping
 - Limit the depth of the tree.
 - Not considering splits that do not reduce classification error
 - Do not split for less data.
- Pruning
 - Bottom up
 - $\text{Cost} = \text{Error}(T) + \lambda * \text{\#of leaves}$




BOOSTING

- Weak learner vs complex model.
- Need to find the sweet spot.
- Two ways:
 - Increase the complexity of the model. (Depth)
 - Boosting!



BOOSTING

- Combines(Boosts) multiple weak learners to increase performance.
- Keep committing mistakes and keep learning from them.
- Algorithm:
 - 1. Learn a weak classifier
 - 2. M = Find the misclassified examples
 - 3. Give higher weights to M
 - 4. Go back to 1



BOOSTING: ADABOOST

- Ensembling: $\text{Pred} = w_1 * h_1 + w_2 * h_2 + w_3 * h_3 + \dots$
- AdaBoost algorithm:
 - Initialize weights $\alpha_i = 1/n$
 - For $t = 1$ to T (where T is the number of weak learners)
 - Learn h_t with data weights α_i
 - Compute coefficient w_t
 - Recompute weights α_i
 - Final model: $\text{pred} = w_1 * h_1 + w_2 * h_2 + w_3 * h_3 + \dots + w_t * h_t$
- Challenges:
 - Finding w_t
 - Finding α_i for each step.

BOOSTING: ADABOOST: FINDING W_T

- Previous error = # of misclassified examples / # of examples.
- New error = weighted error = $(\sum \alpha_i \text{ of misclassified examples}) / (\sum \alpha_i)$
- But how to calculate w_t :
 - w_t is directly proportional to weighted error.
 - $w_t = .5 * \ln((1 - \text{weighted_error}(f_t)) / \text{weighted_error}(f_t))$


BOOSTING: ADABOOST: FINDING α_i

- New weights are inversely proportional to weight.
- $\alpha_i = \alpha_i e^{-w_t}$ (correct)
- $\alpha_i = \alpha_i e^{w_t}$ (incorrect)

BAGGING: RANDOM FOREST

- Ensemble of decision trees
- Wisdom of crowd, poll unrelated models.
- For $t = 1$ to T (T is number of trees in the forest):
 - Select k features randomly.
 - Train h_t
- For prediction: majority vote!

BOOSTING VS BAGGING:



HOMEWORK:

1. Programming:

- a. Extend the decision tree code for multiclass and continuous valued features.
- b. Extend it to implement random forest (and AdaBoost?).

2. Theory:

- a. Study how to do regression using decision trees (Teach us tomorrow morning?)

3. Practical:

- a. Get familiar with decision tree, Random forest, Adaboost, xgboost libraries and use them on some datasets.